CONFLUENT

Clients

Use clients to produce and consume messages

Objectives

Run different producers and consumers against your new cluster.

Use API Keys, Service Users, ACLs and RBAC to secure access to the cluster and its resources.

Labs

Setup

- Create a couple of topics (skip the schema creation)
    - A topic called “test-topic” with 1 partition
    - A topic called “customer” with 6 partitions
        - Have a look at the advanced settings as well
- Create a Java client configuration on the “Clients/Add New Client” page
    - Create an API Key
    - Go directly to the step 4
    - Copy the completed configuration and store it in a local file, for example “ccloud.properties”

Standard Java Clients

The Confluent platform comes with a set of standard Java clients. This part of the lab will ensure you know how to use them.

- On your local machine (laptop), run the following Java clients[1] in separate terminals to see the produced messages in the consumer. Use the “test-topic” or your own topics for this (do not use the “customer” topic).
    - kafka-console-producer
    - kafka-console-consumer

[1] See references for the location of the client tar file, if not installed already

- Remember: You will have to specify the --bootstrap-server again for each command
- Use the "--help" option to find out how to pass the appropriate configuration for producers, consumers and kafka-topics.

Application with schema

The Schema Registry in the Confluent Cloud uses a separate API Key. The following is an example of a client application that communicates with both the broker and the schema registry and therefore requires additional authentication.

- Run a Java application that uses the Schema Registry, for example
  - Download or clone https://github.com/sknop/kafka_transactions
  - Install Maven if you have not done so already (`brew install maven` on a Mac)
  - Compile the applications with "`mvn clean package -DskipTests`"
  - Create a schema API Key and API Secret
    - Navigate to the environment overview for the environment you created ("bootcamp")
    - In the right-hand panel on the bottom, click on the link below Credentials to create a new key
  - Update your config file to add the schema registry

```
# Required connection configs for Confluent Cloud Schema Registry
schema.registry.url=https://psrc-4xrp1.eu-central-1.aws.confluent.cloud
basic.auth.credentials.source=USER_INFO
basic.auth.user.info=<SR-KEY>:<SR-SECRET>
```

  - Run (use option -h to see all options)

```
java -cp target/kafka_transactions-2.0-SNAPSHOT.jar \
  producer.CustomerProducer -c ccloud.properties -v -m 10
```

  - You can use consumer.CustomerConsumer to read the messages (they are Avro encoded, so the normal kafka-console-consumer is not helpful)

Run the kafka-avro-console-consumer

The kafka-avro-console-consumer is an extension of the standard kafka-console-consumer with the Confluent package.

- Run the kafka-avro-console-consumer against your Confluent Cloud instance's “customer” topic. You will need to specify “--from-beginning” to see all messages. You also need a set of additional properties to talk to and authenticate against the schema registry. You can use the “--property” option multiple times. Values to set are
    - schema.registry.url
    - basic.auth.credentials.source=USER_INFO
    - basic.auth.user.info
    - key-deserializer (The key is an Integer)
    - print.key=true
    - key.separator=":"

C/C++ clients - kafkacat/kcat (Optional)

kafkacat (now named kcat) is an example of a C/C++ client based on the same library as the C#/Python/etc clients. On Mac it can be installed with “brew install kcat”

- Run kafkacat/kcat). Use the New Client page and the C/C++ language to create the configuration file, for example, bootcamp-kcat.cfg.
    - Use kcat to read from your topics.

```
kcat -F bootcamp-kcat.cfg  -C -t test-topic
```

- You can use kcat also to read Avro-encoded records:

```
kcat -F bootcamp-kcat.cfg  -C -t customer -s key=i -s value=avro
```

- You will have to add the schema.registry.url with api-key and api-secret to your config-file
    - schema.registry.url=https://api-key:api-secret@URL

Confluent CLI

This section uses the Confluent CLI, which you need to download and install (See references for how).

- Use Confluent CLI to consume from your “test-topic”. There is a “CLI and Tools” section at the left bottom of your Confluent Cloud UI with useful hints, or you can use the Confluent documentation.
- Use the same CLI to consume from “customer”. Now you need to specify serializers since the data is stored in Avro. Try adding the options “`--value-format avro --print-key --key-format integer --delimiter "-"` to see if you can get it to work.

References

https://docs.confluent.io/platform/current/installation/installing_cp/zip-tar.html#install-cp-using-zip-and-tar-archives
https://www.confluent.io/en-gb/blog/using-apache-kafka-command-line-tools-confluent-cloud/
https://docs.confluent.io/confluent-cli/current/install.html#install-confluent-cli
https://docs.confluent.io/confluent-cli/current/beginner-cloud.html
https://developer.confluent.io/tutorials/kafka-console-consumer-producer-avro/confluent.html
https://docs.confluent.io/platform/current/tools/kafkacat-usage.html
(Note that you can download the configuration from your Confluent Cloud cluster / New Client page and use the option -F with that file, rather than passing all parameters in on the command line).

Expected Outcomes

Access Confluent Cloud cluster and produce to and consume from topics using different clients.
Understand how API Keys are used for authentication.

Check your understanding

This colour marks advanced questions.

- What is an API Key?
- How and why do the configurations for Java and C/C++ clients differ?
- Why are the API Keys for the Kafka cluster and the Schema Registry different?